

Optimising Llama2

Enhancing Efficiency and Scalability of Large Language Models with
PyTorch

Chris Stylianou
Research Engineer
CaSToRC
eurocc.cyi.ac.cy





A Bit About Myself

- Research Engineer at CaSToRC
- Collaborations Task Leader for EuroCC2
- Area of expertise in High Performance Computing (HPC)
- PhD Candidate, EPCC
- Contact & Info:
 - Email: c.stylianou@cyi.ac.cy
 - Website: cstyl.github.io





- **Generative AI** use cases have **exploded** in popularity recently!
- **Text generation** one particularly popular area.
 - ChatGPT, Llama, vLLM etc.
- PyTorch is a popular **open-source** ML/AI library widely used in the area of AI
 - **Ease of Use** and Flexibility
 - Strong community and industry support
 - **Integration** with Python ecosystem
 - Accelerated computing (**GPU support**)
 - Educational resources
- How fast we can run transformer **inference** with only pure, native PyTorch?



- Training GitHub Repo to be available at:
 - <https://github.com/CaSToRC-Cyl/isc2024-tutorial/tree/main/gpt-fast>
 - Follow README .md for every part in the slides.
- Complete tutorial:
 - <https://github.com/pytorch-labs/gpt-fast>
- To download Llama2 7B parameters model, go to
 - <https://huggingface.co/meta-llama/Llama-2-7b>
 - Stored in HF Transformer format
 - Shared directory with processed weights available for the training!
- Hardware needed: GPUs with BF16 & INT8 precision
 - E.g., A100 NVIDIA GPUs 40GB



Environment Setup

```
# Connect to GWDG
$ ssh -i /path/to/ssh-key <username>@glogin.hlrn.de
# Shared Project Path
$ export SHARED_PATH=/mnt/lustre-emmy-ssd/projects/isc2024_accel_genai_pytorch
# Make a local copy of the code
$ git clone https://github.com/CaSToRC-CyI/isc2024-tutorial.git
# Activate environment
$ module load python/3.9.16
$ source $SHARED_PATH/torch/bin/activate
# Code for the session
$ cd isc2024-tutorial/gpt-fast
```

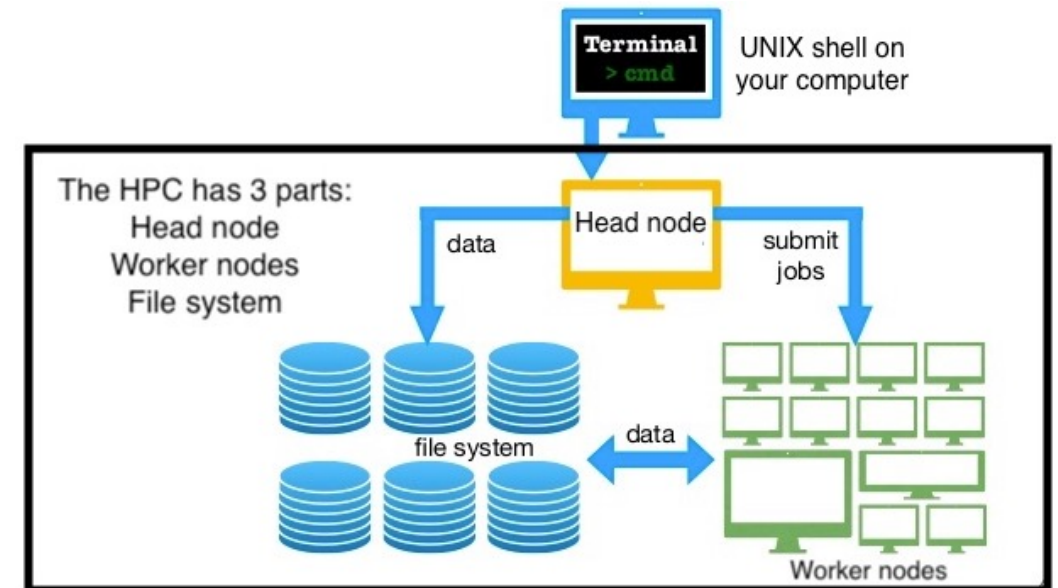


Figure 1: High-level overview of a Supercomputer [1].



A Note on the Scheduler Settings

```
#!/bin/bash -l

#SBATCH --job-name=gpt-fast-baseline
#SBATCH --nodes=1           # Request 1 Node
#SBATCH --ntasks-per-node=1 # Request 1 Process
#SBATCH -G A100:1           # Request 1 GPU
#SBATCH --cpus-per-task=16  # Request 12 threads (not used, but good for locality)
#SBATCH --partition=grete   # Request on Cluster with A100 NVIDIA GPUs
#SBATCH --time=01:00:00     # Request access for 1 hour
#SBATCH --exclusive         # Request exclusive access on whole node
#SBATCH --output=%x.out

# Setup paths and Environment
SHARED_PATH=/mnt/lustre-emmy-ssd/projects/isc2024_accel_genai_pytorch
EXEC_PATH=$SHARED_PATH/isc2024-tutorial/gpt-fast # Path to Code
CHECKPOINTS_PATH=$SHARED_PATH/gpt-fast-checkpoints # Path to Weights

module load python/3.9.16
source $SHARED_PATH/torch/bin/activate

export DEVICE=cuda
export MODEL_REPO=meta-llama/Llama-2-7b
python $EXEC_PATH/generate.py --checkpoint_path $CHECKPOINTS_PATH/$MODEL_REPO/model.pth --prompt "Hello, my name is"
```



- **Llama2** (Meta) model
 - **7B** Parameters (~13GB)
 - **BF16** Precision used for weights
- Inference
 - Prompt: “Hello, my name is”
 - Generate **5 samples**, up to **200 tokens**.
 - **Batch Size = 1**
- Resources configuration:
 - # of GPUs (Nodes): 1 (1)
 - CUDA Version: 11.8.0
 - PyTorch Version: 2.3

“Hello, my name is [Name], and I am a [insert activity or hobby here] enthusiast. I have been involved in this activity for [insert number of years] years, and I can tell you that it has brought me a great deal of joy and fulfillment.”


Original Model:

<https://github.com/meta-llama/llama/tree/main>

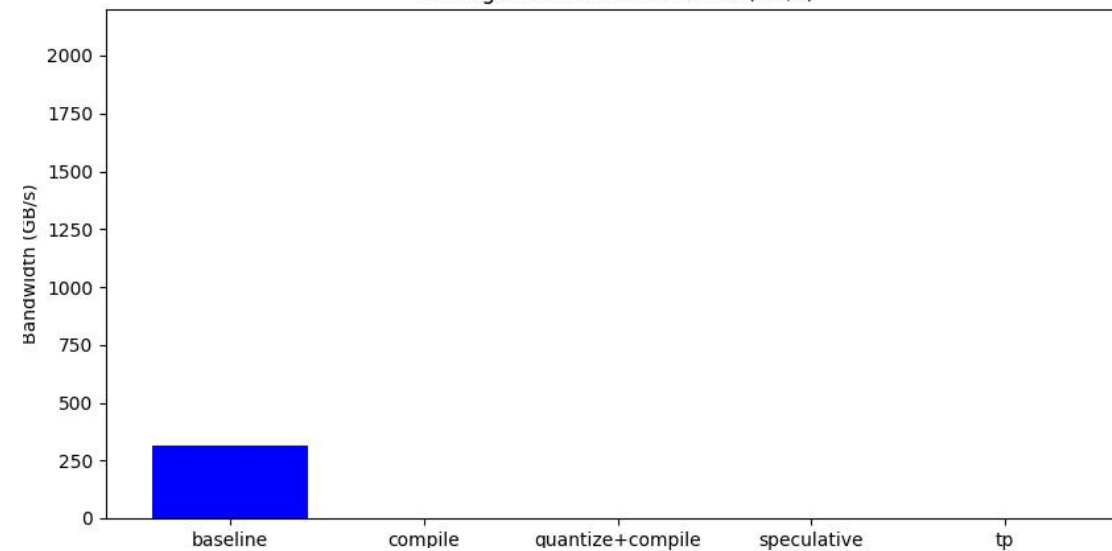


Running Baseline Version (23.17 tokens/s)

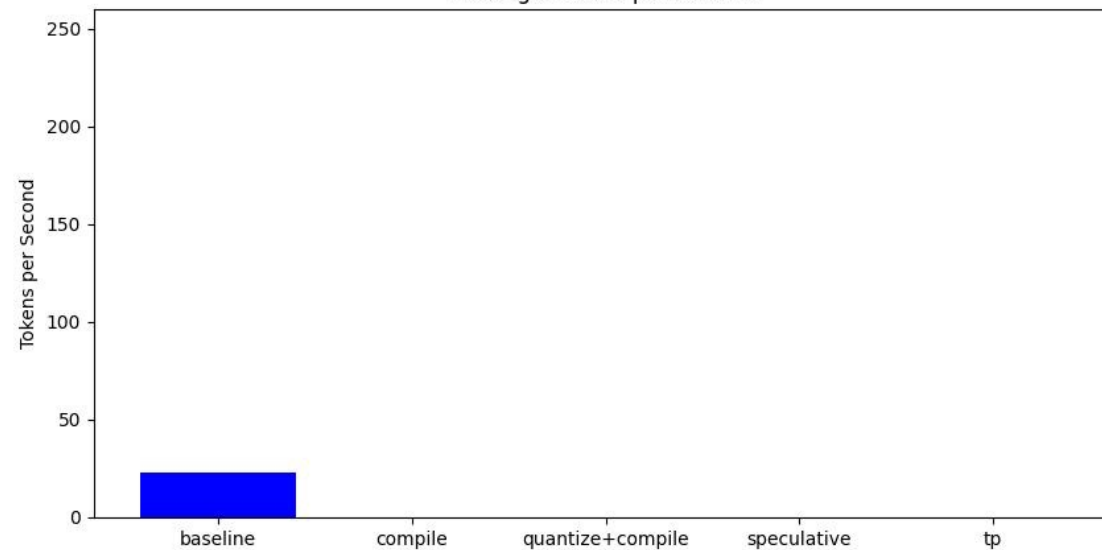
Memory Used (GB)



Average Bandwidth Achieved (GB/s)



Average Tokens per Second





- Modern GPUs are **extremely efficient**, with a lot compute power.
- Deep learning computations run **entirely on GPU**
 - CPU acts as the **orchestrator**, telling the GPU what to do next.
- Complex algorithms often require **multiple GPU operations**
→ **Multiple kernel launches**
- Each kernel launch is associated with **some overhead**.
- When **kernels** sent **individually** and **complete quickly**:
 - CPU tries to keep up with the GPU (**CPU-bound**)
 - **Combined overhead** becomes noticeable, drop in overall performance!
- Solutions:
 - Rewrite in C++, raw CUDA?
 - Send more work to GPU at once?



- **Purpose of `torch.compile()`:** [Requires PyTorch ≥ 2.0]
 - **Arbitrary Python functions** can be optimised by passing the callable to `torch.compile()`
 - JIT Compilation, via minimal code changes.
- **Usage Example:**

```
def foo(x, y):  
    a = torch.sin(x)  
    b = torch.cos(y)  
    return a + b  
  
opt_fool = torch.compile(foo)  
print(opt_fool(torch.randn(10, 10), torch.randn(10, 10)))
```
- **Key Considerations:**
 - Can also be used for `torch.nn.Module`
 - Compiles the model into optimised kernels during execution!
 - Useful if optimised model is used several times (Only have to pay once!)



- Speedup achieved via reducing **Python overhead** and **GPU read/writes**.
- `torch.compile` optimizes computational graphs in PyTorch to **reduce CPU overhead**.
 - Allows for larger model components to be **compiled into a single optimized section**.
- **Usage Example:**

```
torch.compile(decode_one_token, mode="reduce-overhead",  
              fullgraph=True)
```
- **Available Modes:**
 - *Default*: For large models, **low compile time**, no extra memory
 - *Reduce-overhead*: **Reduces framework overhead**, uses extra memory, good for **small batches**
 - *Max-autotune*: Produces the fastest model but takes a **very long time to compile**.
- `fullgraph=True` :
 - Minimizes the frequency and impact of "graph breaks"
 - i.e portions that cannot be compiled.
 - Ensures maximum potential utilization of `torch.compile`.



- A common optimisation trick for speeding up **transformer inference**.
 - Activations computed for the previous tokens are cached.
- As more tokens are generated, the “logical length” of the kv-cache grows.
 - Rellocating and copying every time the cache grows!
- Due to this dynamism `torch.compile` **less efficient**
 - i.e. need to recompile.



- Use “**static**” kv-cache:

- Statically allocate the maximum size of kv-cache
- **Mask out the unused values** in the attention portion

```
with torch.device(device):  
    model.setup_caches(max_batch_size=1, max_seq_length=max_seq_length)
```

- Two-phase compilation:

1. The prefill, where the **entire prompt is processed**

- more **dynamic**, due to variable prompt length

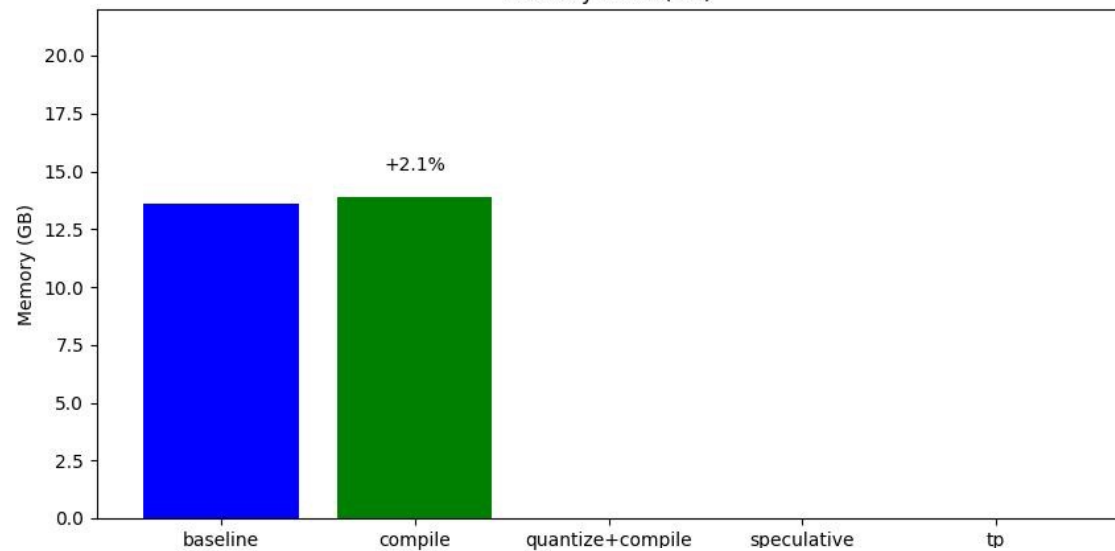
```
prefill = torch.compile(prefill, dynamic=True,  
    fullgraph=True)
```

2. Decoding, where each token is generated (kv-cache static)

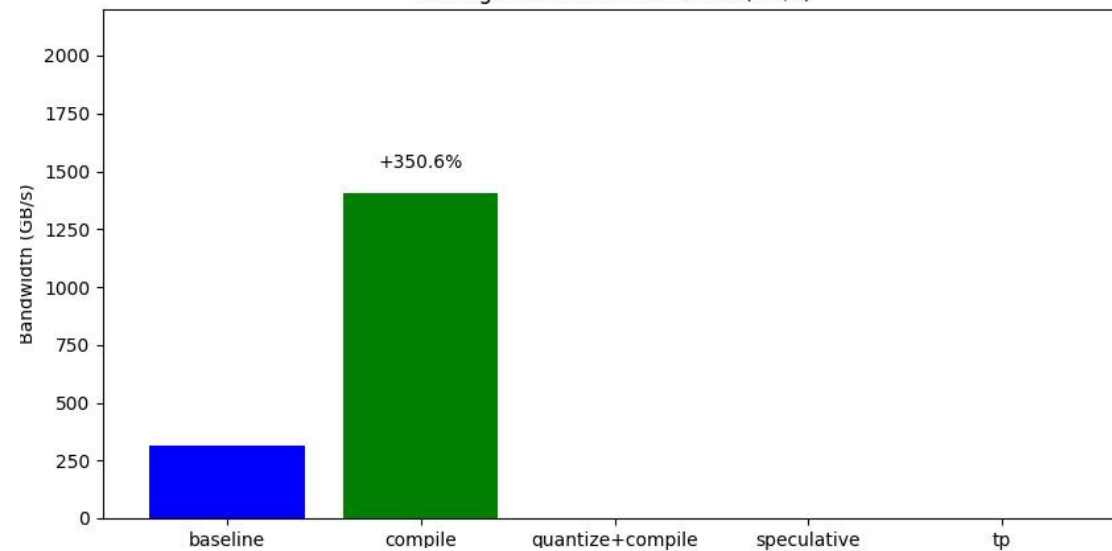
```
decode_one_token = torch.compile(decode_one_token,  
    mode="reduce-overhead",  
    fullgraph=True)
```




Memory Used (GB)



Average Bandwidth Achieved (GB/s)



Average Tokens per Second





- Largest bottleneck now is the cost of loading weights **from GPU global memory to registers**.
 - Happens in each forward pass.
- **Percentage of memory bandwidth** used during inference given by **Model Bandwidth Utilisation (MBU)**:

$$MBU = \frac{\# Params * \frac{bytes}{param} * \frac{tokens}{second}}{Memory Bandwidth}$$

- In our case, (7B Params, each FP16, and 104 tokens/s):

$$MBU = \frac{7B * 2 * 104}{1.6 TB} = 93\%$$

- Very close to the theoretical limit!
- We can change how many **bytes each parameter is stored in!**



- Quantize only the weights, **computation still done in BF16**.
 - Easy to apply with little to no loss of accuracy.
- Done once, **offline!**
 - Results in reduced memory footprint and faster execution on hardware.
- Here applied per-channel.
- Example:

$$X_q = \text{round} \left(\frac{127}{\max|X|} * X \right)$$

maps FP16 values into [-127, 127] 8-bit integers!


- Quantized Matrix Multiplication becomes:

```
x: bf16[1, K]
weight: int8[K, N]
@torch.compile
def int8_mm(x, weight):
    return F.linear(x, weight.to(torch.bfloat16))
```




Running Quantized Version (133.07 tokens/s)


Memory Used (GB)



Average Bandwidth Achieved (GB/s)



Average Tokens per Second



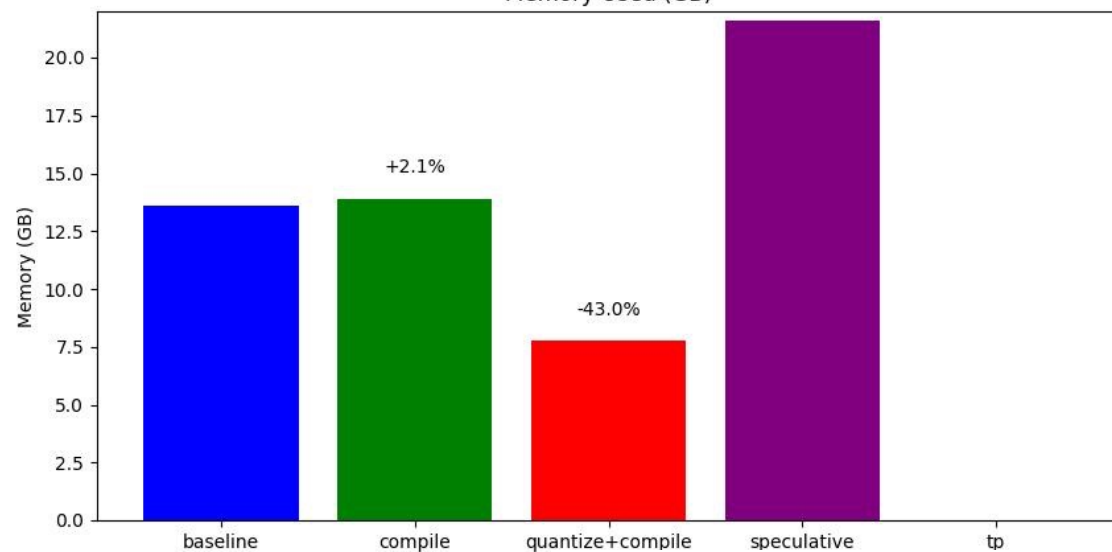


- For every token generated, weights have to be loaded over and over.
 - **Strict serial dependency** in autoregressive generation.
- **Speculative decoding** breaks this dependency!
- Main Idea:
 - Larger model, which we want to use for inference (**Verifier Model**)
 - Smaller model, able to generate text much faster (**Draft Model**)
 - But less accurate!
 - Generate **N tokens using the cheaper** draft model, then process all of them in parallel using the verifier model
 - Those **not matching**, discard and **regenerate with Verifier Model**.
- Speculative decoding **does not change the quality** of the output.
- Around 50 lines of code implementation.
- **Runtime performance varies** depending on the generated text.

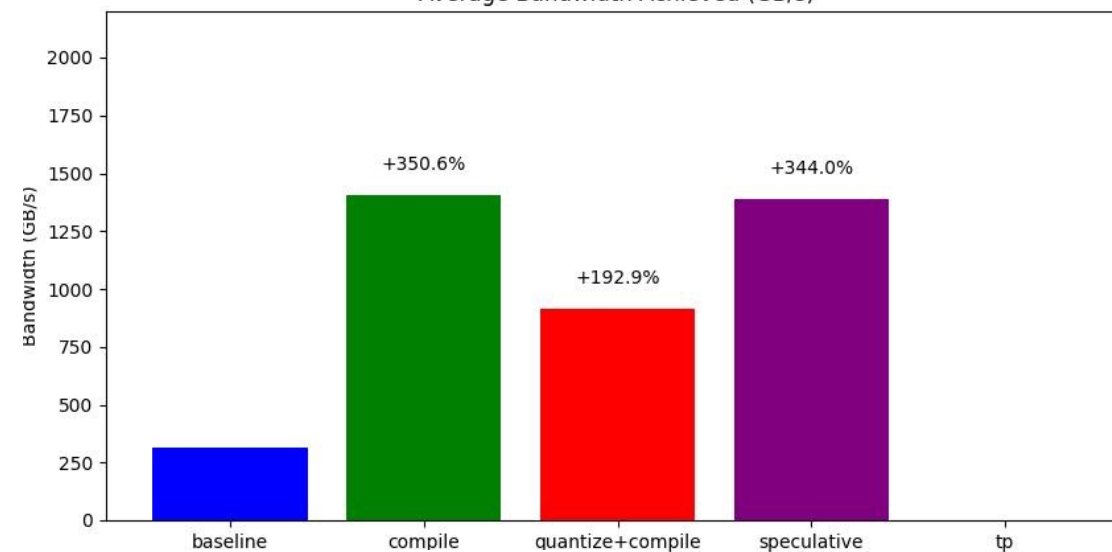


Running Speculative Version (102.85 tokens/s)


Memory Used (GB) +58.6%



Average Bandwidth Achieved (GB/s)



Average Tokens per Second








- So far, only one GPU was used!
- Running on more GPUs gives us access to more memory bandwidth.
- Parallelisation strategy is to **split the processing of one token across multiple devices**
 - Tensor Parallelism
- PyTorch supports this, although low-level API currently.
 - 150 lines of code, no model changes!
- Main idea:

```
def forward(self, x: torch.Tensor) -> torch.Tensor:
    x = F.silu(self.c_fc1(x)) * self.c_fc2(x)
    x = self.c_proj(x)
    x = collectives.all_reduce(x, "sum",
                               list(range(LOCAL_WORLD_SIZE)))
    return x
```




Running Tensor Parallelism Version (250 tokens/s)





- Memory efficient attention implementations
 - Scaled dot product attention (SPDA)
- Semi-structured (2:4) Sparsity
 - Good for Sparsification/Pruning



From developer.nvidia.com/blog/exploiting-ampere-structured-sparsity-with-cusparselt



- Using native PyTorch offers ease of use without sacrificing performance.
- The code for optimisations is around 900 lines.
 - torch.compile
 - Quantization (BF16 to INT8)
 - Speculative Decoding (Expert 7B 16BF, Draft 7B INT8)
 - Tensor Parallelism
- From:
 - **23.17 tokens/second to 250!**
 - 1 GPU to 4 GPUS.
 - Overhead-bound to memory bandwidth bound.



This project has received funding from the European High-Performance Computing Joint Undertaking (JU) under grant agreement No 101101903.





1. <https://gu-ereseach.github.io/hpcWorkshop/content/12-logOntoHPC.html>